

Selected Exercises from Fundamentals of Database Systems

Andrew Kowalczyk

Chapter 1

1.8 Identify some informal queries and update operations that you would expect to apply to the database shown in Figure 1.2.

Queries

1. Find the names of all students majoring in Mathematics.
2. What are the prerequisites of the Database course?
3. Find the transcript of the student named Brown. We would find <Course_name, Section_identifier, Semester, Year, Grade> for each course section that Brown has taken.

Updates

1. Insert a new student in the database whose Name=Kowalczyk, Student_number=25, Class=4, and Major=CS.
2. Change the grade that Brown received in Discrete Mathematics to a D.

1.9 What is the difference between controlled and uncontrolled redundancy? Illustrate with examples.

Redundancy is the term given when the same data is stored multiple times in several places in a database. If you look at Figure 1.5(a) in the text, you can see that the name of the student with Student_number=8 is Brown is stored multiple times. Redundancy is controlled when the database management system (DBMS) ensures that multiple copies of the same data are consistent. To illustrate this, let's say we are adding a new record with Student_number=8 to be stored in the database of Figure 1.5(a). If we were to have uncontrolled redundancy, the DBMS would have no control over this. If we were to have controlled redundancy, the DBMS would ensure that Student_name=Brown in that record.

1.10 Specify all the relationships among the records of the database shown in Figure 1.2.

1. Every GRADE_REPORT record is related to one STUDENT record and one SECTION record.
2. Every SECTION record is related to a COURSE record.
3. Every PREREQUISITE record relates two COURSE records. One being a course and the other being a prerequisite to that course.

1.11 Give some additional views that may be needed by other user groups for the database shown in Figure 1.2.

1. A view of each class section that groups all the students who took that section and their respective grade.
2. A view that gives the number of courses taken and the grade point average for each student.

1.12 Cite some examples of integrity constraints that you think can apply to the database shown in Figure 1.2.

Key constraints

1. Student_number must be unique for each STUDENT record.
2. Course_number must be unique for each COURSE record.

Referential integrity constraints

1. The Course_number in a SECTION record must also exist in some COURSE record.
2. The Student_number in a GRADE_REPORT record must also exist in some STUDENT record.

Domain constraints

1. Grades in a given GRADE_REPORT record must be one of these values: A, B, C, D, F, I, W .

Chapter 2

2.12 Think of different users for the database shown in Figure 1.2. What types of applications would each user need? To which user category would each belong, and what type of interface would each need?

1. Students add and drop classes. Actions that they can do are as listed:
 - (a) Register themselves in a section of a course
 - (b) Drop themselves from a section of a course
2. Registrar. They enter data of registration of students in sections of courses, and later enter the grades of the students. Actions that they can do are as listed:
 - (a) Check whether a student who is registered in a course has the appropriate prerequisite courses
 - (b) Add a student to a section of a course
 - (c) Enter the student grades for a section
3. Admissions. Their main application would be to enter newly accepted students into the database. Actions that they can do are as listed:
 - (a) Add students to the school's records

Chapter 3

Chapter 4

Chapter 5

Chapter 6

6.15. Show the result of each of the sample queries in Section 6.5 as it would apply to the database state in Figure 3.6.

Query 1: Find the name and address of all employees who work for the 'Research' department.

FNAME	LNAME	ADDRESS
John	Smith	731 Fondren, Houston, TX
Franklin	Wong	638 Voss, Houston, TX
Ramesh	Narayan	975 Fire Oak, Humble, TX
Joyce	English	5631 Rice, Houston, TX

Query 2: For every project located in 'Stafford', list the project number, the controlling department number, and the department manager's last name, address, and birth date.

PNUMBER	DNUM	LNAME	ADDRESS	BDATE
10	4	Wallace	291 Berry, Bellaire, TX	20-JUN-31
30	4	Wallace	291 Berry, Bellaire, TX	20-JUN-31

Query 3: Find the names of all employees who work on all the projects controlled by department number 5.

LNAME	FNAME
-------	-------

Query 4: Make a list of project numbers for projects that involve an employee whose last name is 'Smith' as a worker or as a manager of the department that controls the project.

PNO
1
2

Query 5: List the names of all employees with two or more dependents.

LNAME	FNAME
Smith	John
Wong	Franklin

Query 6 List the names of employees who have no dependents.

LNAME	FNAME
Zelaya	Alicia
Narayan	Ramesh
English	Joyce
Jabbar	Ahmad
Borg	James

Query 7: List the names of managers who have at least one dependent.

LNAME	FNAME
Wallace	Jennifer
Wong	Franklin

6.16 Specify the following queries on the COMPANY relational database schema shown in Figure 5.5, using the relational operators discussed in this chapter. Also show the result of each query as it would apply to the database state in Figure 3.6.

I use the symbol σ for SELECT, π for PROJECT, \bowtie for EQUIJOIN, $*$ for NATURAL JOIN, and \mathfrak{F} for FUNCTION.

(a) Retrieve the names of employees in department 5 who work more than 10 hours per week on the "ProductX" project.

Relational Operators

$EMP_W_X \leftarrow (\sigma_{PNAME='ProductX'}(PROJECT)) \bowtie_{(PNUMBER),(PNO)} (WORKS_ON)$
 $EMP_WORK_10 \leftarrow (EMPLOYEE) \bowtie_{(SSN),(ESSN)} (\sigma_{HOURS>10}(EMP_W_X))$
 $RESULT \leftarrow \pi_{LNAME,FNAME} (\sigma_{DNO=5}(EMP_WORK_10))$

Result

LNAME	FNAME
Smith	John
English	Joyce

(b) List the names of employees who have a dependent with the same first name as themselves.

Relational Operators

$E \leftarrow (EMPLOYEE) \bowtie_{(SSN,FNAME),(ESSN,DEPENDENT_NAME)} (DEPENDENT)$
 $R \leftarrow \pi_{LNAME,FNAME} (E)$

Result

LNAME	FNAME
-------	-------

(c) Find the names of employees that are directly supervised by "Franklin Wong".

Relational Operators

$WONG_SSN \leftarrow \pi_{SSN} (\sigma_{FNAME='Franklin' \text{ AND } LNAME='Wong'}(EMPLOYEE))$
 $WONG_EMPS \leftarrow (EMPLOYEE) \bowtie_{(SUPERSSN),(SSN)} (WONG_SSN)$
 $RESULT \leftarrow \pi_{LNAME,FNAME} (WONG_EMPS)$

Result

LNAME	FNAME
Smith	John
Narayan	Ramesh
English	Joyce

(d) For each project, list the project name and the total hours per week (by all employees) spent on that project.

Relational Operators

$\text{PROJ_HOURS}(\text{PNO}, \text{TOT_HRS}) \leftarrow \text{PNO} \bowtie \text{SUM HOURS} (\text{WORKS_ON})$
 $\text{RESULT} \leftarrow \pi_{\text{PNAME}, \text{TOT_HRS}} ((\text{PROJ_HOURS}) \bowtie_{(\text{PNO}), (\text{PNUMBER})} (\text{PROJECT}))$

Result

PNAME	TOT_HRS
ProductX	52.5
ProductY	37.5
ProductZ	50.0
Computerization	55.0
Reorganization	25.0
Newbenefits	55.0

(e) Retrieve the names of employees who work on every project.

Relational Operators

$\text{PROJ_EMPS}(\text{PNO}, \text{SSN}) \leftarrow \pi_{\text{PNO}, \text{ESSN}} (\text{WORKS_ON})$
 $\text{ALL_PROJS}(\text{PNO}) \leftarrow \pi_{\text{PNUMBER}} (\text{PROJECT})$
 $\text{EMPS_ALL_PROJS} \leftarrow \text{PROJ_EMPS} \div \text{ALLPROJS}$
 $\text{RESULT} \leftarrow \pi_{\text{LNAME}, \text{FNAME}} (\text{EMPLOYEE} * \text{EMP_ALL_PROJS})$

Result

LNAME	FNAME
-------	-------

(f) Retrieve the names of employees who do not work on any project.

Relational Operators

$\text{ALL_EMPS} \leftarrow \pi_{\text{SSN}} (\text{EMPLOYEE})$
 $\text{WORKING_EMPS}(\text{SSN}) \leftarrow \pi_{\text{ESSN}} (\text{WORKS_ON})$

$\text{NON_WORKING_EMPS} \Leftarrow \text{ALL_EMPS} - \text{WORKING_EMPS}$
 $\text{RESULT} \Leftarrow \pi_{\text{LNAME}, \text{FNAME}} (\text{EMPLOYEE} * \text{NON_WORKING_EMPS})$

Result

LNAME	FNAME
-------	-------

(g) For each department, retrieve the department name and the average salary of all employees working in that department.

Relational Operators

$\text{DEPT_AVG_SALS}(\text{DNUMBER}, \text{AVG_SAL}) \Leftarrow \text{DNO} \Join \text{AVG_SALARY} (\text{EMPLOYEE})$
 $\text{RESULT} \Leftarrow \pi_{\text{DNUMBER}, \text{AVG_SAL}} (\text{DEPT_AVG_SALS} * \text{DEPARTMENT})$

Result

DNUMBER	AVG_SAL
Research	33250
Administration	31000
Headquarters	55000

(h) Retrieve the average salary of all female employees.

Relational Operators

$\text{RESULT}(\text{AVG_F_SAL}) \Leftarrow \Join \text{AVG_SALARY} (\sigma_{\text{SEX}='F'} (\text{EMPLOYEE}))$

Result

AVG_F_SAL
31,000

(i) Find the names and addresses of employees who work on at least one project located in Houston but whose department has no location in Houston.

Relational Operators

$\text{E_P_HOU}(\text{SSN}) \Leftarrow \pi_{\text{ESSN}} (\text{WORKS_ON} \bowtie (\text{PNO}), (\text{PNUMBER}) (\sigma_{\text{PLOCATION}='Houston'} (\text{PROJECT})))$
 $\text{D_NO_HOU} \Leftarrow \pi_{\text{DNUMBER}} (\text{DEPARTMENT}) - \pi_{\text{DNUMBER}} (\sigma_{\text{DLOCATION}='Houston'} (\text{DEPARTMENT}))$
 $\text{E_D_NO_HOU} \Leftarrow \pi_{\text{SSN}} (\text{EMPLOYEE} \bowtie (\text{PNO}), (\text{DNUMBER}) (\text{D_NO_HOU}))$
 $\text{RESULT_EMPS} \Leftarrow \text{E_P_HOU} - \text{E_D_NO_HOU}$
 $\text{RESULT} \Leftarrow \pi_{\text{LNAME}, \text{FNAME}, \text{ADDRESS}} (\text{EMPLOYEE} * \text{RESULT_EMPS})$

Result

LNAME	FNAME	ADDRESS
Wallace	Jennifer	291 Berry, Bellaire, TX

(j) List the last names of department managers who have no dependents.

Relational Operators

$\text{DEPT_MANAGERS}(\text{SSN}) \leftarrow \pi_{\text{MGRSSN}} (\text{DEPARTMENT})$
 $\text{EMPS_WITH_DEPENDENTS}(\text{SSN}) \leftarrow \pi_{\text{ESSN}} (\text{DEPENDENT})$
 $\text{RESULT_EMPS} \leftarrow \text{DEPT_MANAGERS} - \text{EMPS_WITH_DEPENDENTS}$
 $\text{RESULT} \leftarrow \pi_{\text{LNAME}, \text{FNAME}} (\text{EMPLOYEE} * \text{RESULT_EMPS})$

Relational Operators

LNAME	FNAME
Borg	James

Chapter 7

Chapter 8

Chapter 9

Chapter 10